




Figure 1: Aggregated normalized best evaluation returns across 10 MuJoCo tasks. For each task, the return is normalized to $[0, 1]$ using the minimum and maximum returns across all methods. The interquartile mean (IQM) excludes the top and bottom two task scores for each method to mitigate the effect of outliers.




Figure 2: Aggregated normalized best evaluation returns across 8 DeepMind Control Suite tasks. For each task, the return is normalized to $[0, 1]$ by dividing by 1000.